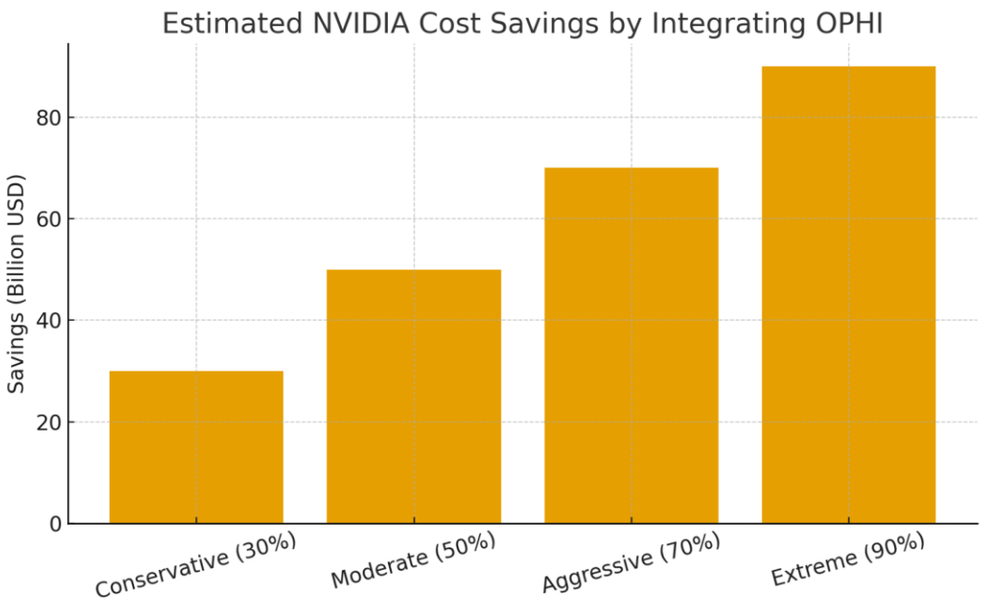

OPHI vs GPU-Centric Spend — Savings Snapshot

- Baseline: $100B GPU-centric investment (chips, 10GW DCs, power/cooling, training, ops, equity).
- What OPHI obsoletes: brute-force matrix burns, heavy checkpointing, scraping-driven data ops.
- Quality gates: coherence/entropy/drift → fewer retries, smaller artifacts, built-in provenance.
- Scenario savings vs $100B: 30%=$30B, 50%=$50B, 70%=$70B, 90%=$90B.
- Component model (assumed): silicon 70%, DC 60%, energy 80%, ops 50%, compliance 90% reductions.
- Plausible range: $50–$70B avoided; midpoint reference: ~$60B.
- Caveat: Assumptions illustrative; adjust shares/reductions in the CSV to stress-test.

Generated: 2025-09-24T00:54:32.848192+00:00